

‘More Bikes’: Experiments in Univariate Regression

Tim Lawson

University of Bristol

tim.lawson@bristol.ac.uk

1 Task description

The assignment is to predict the number of available bikes at 75 rental stations in three hours’ time for a period of three months from November 2014, i.e., a supervised univariate regression problem. It is divided into three sub-tasks, which differ in the information that is available. For sub-task 1, the number of available bikes at each of the 75 stations for the month of October 2014 is provided (section 4). This sub-task may be approached by building a separate model for each station or a single model for all 75 stations. For sub-task 2, a set of linear models that were trained on the number of available bikes at each of a separate set of 200 stations are provided (section 5). Finally, for sub-task 3, both sources of information may be used. Additionally, the number of available bikes at the first ten stations between June 2012 and October 2014 is provided for analysis (see fig. 3).

The predictions are evaluated by the mean absolute error (MAE) between the predicted and true numbers of available bikes. Hereafter, the MAE is referred to as the ‘score’.

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i| \quad (1)$$

The evaluation data was not available to participants, but the score achieved on a held-out test set is reported on the task leaderboard.¹ This report begins with a description of the general approach taken to the assignment in section 2, and section 3 presents preliminary data analysis. Finally, sections 4 and 5 detail the results of the model classes that I applied to each of the sub-tasks.² I give the score achieved on the held-out test set by the best estimators for each model class alongside the mean scores on cross-validation folds of the data provided for sub-task 1 in table 4.

2 Approach

I used the `scikit-learn` Python package (Pedregosa et al. 2011) throughout this report. In each case, preprocessing and feature selection were performed by *estimators* that implemented the *transformer* interface, prediction was performed by estimators that implemented the *predictor* interface, and estimators were composed into Pipeline objects over which hyperparameter search was performed (Buitinck et al. 2013, pp. 4–9).

Generally, standard k -fold cross-validation is disfavoured for time-series data due to the inherent correlation between successive folds (Bergmeir et al. 2018). Instead, I used nested time-series cross-validation³ with ten folds to evaluate the models, which is illustrated in fig. 1. I determined the best estimator for each model class by grid search⁴ with the mean absolute error as the scoring function,

¹<https://www.kaggle.com/c/morebikes2023/leaderboard>

²The code that produced these results is available at <https://github.com/tslwn/more-bikes>.

³See `sklearn.model_selection.TimeSeriesSplit`.

⁴See `sklearn.model_selection.GridSearchCV` and `sklearn.model_selection.HalvingGridSearchCV`.

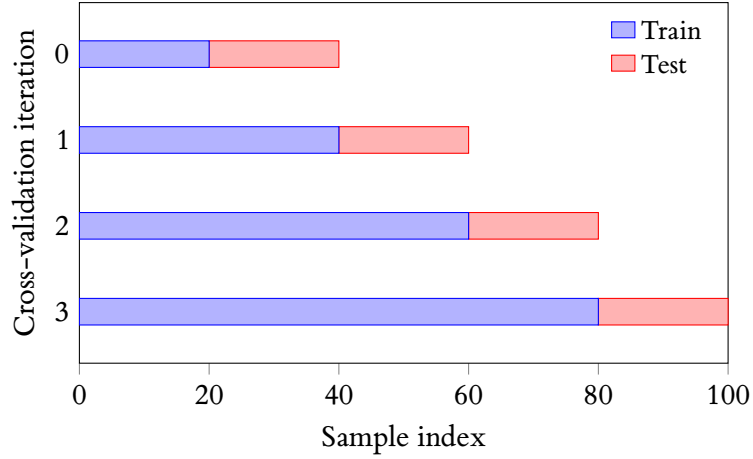


Figure 1: A visualization of the nested time-series cross-validation behaviour that I used, after the documentation for `sklearn.model_selection.TimeSeriesSplit`.

following the task description (section 1). Finally, I assessed the statistical significances of the differences between the mean scores of the best estimators by paired t -tests of the scores on the cross-validation folds. I describe these methods in more detail in sections 4 and 5.

3 Data analysis

The data provided for sub-task 1 is recorded at hourly intervals with $n = 54385$ instances across the 75 stations. A summary of its features and the distributional characteristics of the non-temporal quantitative features are given in tables 1 and 2. This analysis revealed that many of the features and the target variable bikes are missing for $n = 73$ instances, which were henceforth excluded. Additionally, the ‘profile’ features, i.e., the features derived from the numbers of available bikes at preceding times, are not defined for the first week of instances at each station. Hence, they are missing for approximately $\frac{1}{4}$ of the instances. The meteorological features are constant for all stations at a given timestamp. Finally, some features have a natural range – in particular, the number of available bikes at a station is bounded by zero and its number of docks (section 3.2).

3.1 Feature selection

Intuitively, features that have zero variance are not informative for regression, so I automatically excluded them.⁵ In the case of sub-task 1, the available data is limited to the month of October 2014, so these included the month and year. The ‘station’ features (table 1) are constant for all instances at a given station, so were likewise excluded in this case. Finally, the precipitation feature is zero for all instances. Correlations between features are undesirable in regression analysis (Alin 2010), so I also sought to identify and exclude redundant features. Therefore, I computed the Pearson correlation coefficients between pairs of quantitative features (fig. 2), which yielded the following:

- `bikes_avg_full` and `bikes_avg_short` are perfectly correlated ($r = 1.00$).
- `bikes_3h_diff_avg_full` and `bikes_3h_diff_avg_short` are perfectly correlated ($r = 1.00$).
- `wind_speed_max` and `wind_speed_avg` are highly correlated ($r = 0.96$).

Hence, for sub-task 1, the second of each of these pairs of features was also excluded. For sub-task 2, the ‘profile’ features were retained because they are inputs to the pre-trained linear models (table 6).

⁵See `sklearn.feature_selection.VarianceThreshold`.

Category	Feature	Data type	Kind
Station	station	int	ordinal
	latitude	float	
	longitude	float	
	docks	int	
Temporal	timestamp	int	ordinal
	year	int	
	month	int	
	day	int	
	hour	int	
	weekday	str	
	weekhour	int	
	is_holiday	bool	categorical
Meteorological	wind_speed_max	float	
	wind_speed_avg	float	
	wind_direction	float	
	temperature	float	
	humidity	float	
	pressure	float	
	precipitation	float	
Bikes	bikes	int	
	bikes_avg_full	float	
	bikes_avg_short	float	
	bikes_3h	int	
	bikes_3h_diff_avg_full	float	
	bikes_3h_diff_avg_short	float	

Table 1: A summary of the features and the target variable (bikes). Except where indicated, the features are quantitative. The features have been renamed to be easier to read.

Feature	n/a	μ	σ^2
wind_speed_avg	73	4.69	21
wind_direction	365	170.23	7,553.8
temperature	73	21.71	10.7
humidity	73	65.94	279.7
pressure	73	1,002.26	1,808.27
bikes	73	7.35	43.15
bikes_avg_full	12,264	7.31	35.82
bikes_avg_short	12,264	7.31	35.82
bikes_3h	292	7.34	43.17
bikes_3h_diff_avg_full	12,483	$4.1 \cdot 10^{-3}$	22.4
bikes_3h_diff_avg_short	12,483	$4.1 \cdot 10^{-3}$	22.4

Table 2: The numbers of missing values (n/a) and distributional characteristics of the non-temporal quantitative features and the target variable (bikes). Features that have zero variance for the first case of sub-task 1 are excluded (section 3.1).

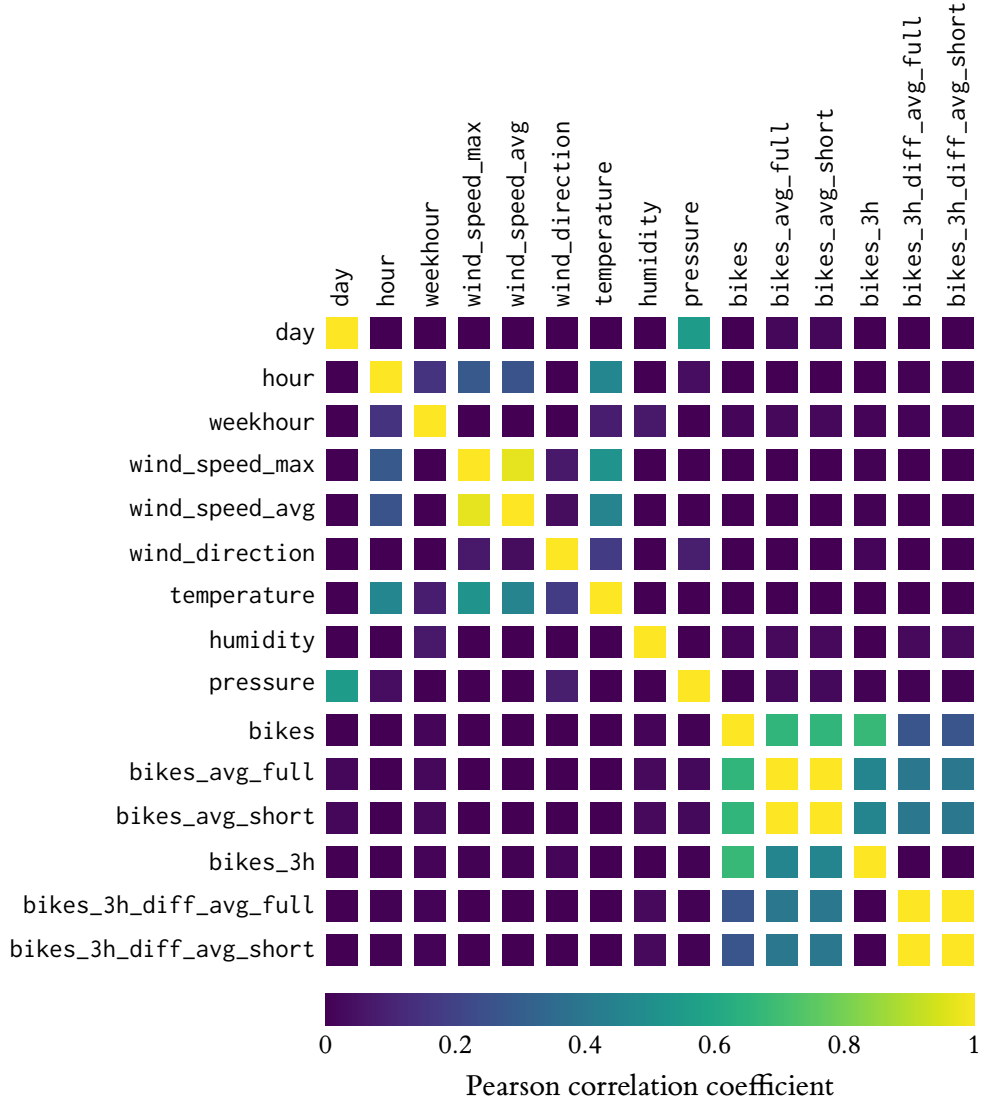


Figure 2: The Pearson correlation coefficients between pairs of the quantitative features and the target variable (bikes). The ordering follows table 1. Features that have zero variance for the first case of sub-task 1 are excluded (section 3.1). The timestamp feature is also excluded because it is naturally correlated with the other temporal features.

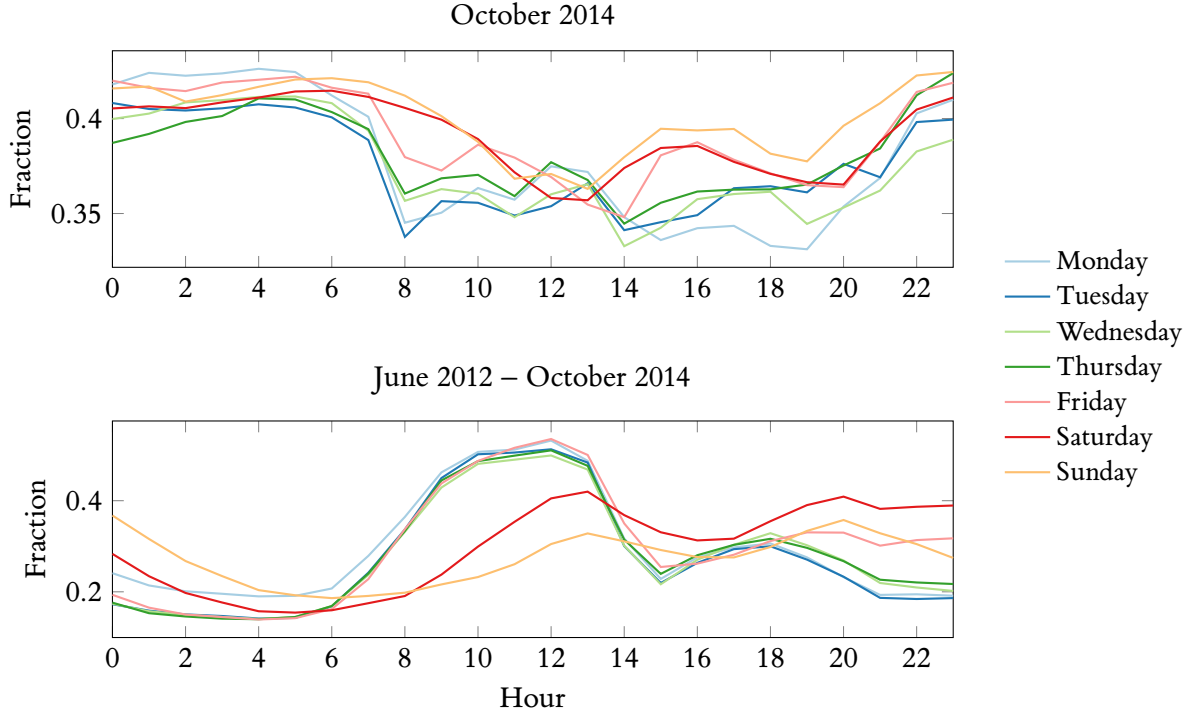


Figure 3: The average fraction of available bikes at each hour of the day, separated by the day of the week. In the top chart, the data is limited to the month of October 2014 for the 75 stations to predict. In the bottom chart, the data from the first ten stations is also included (section 1).

3.2 Feature engineering

For sub-task 1, the bounds on the number of available bikes at each station were enforced by predicting the *fraction* of bikes, i.e., the number of bikes divided by the number of docks. I implemented this by extending the `TransformedTargetRegressor` meta-estimator to permit data-dependent transforms.⁶ Another possible transformation of the target variable would have been to predict the *change* in the number of available bikes, but I did not explore this option.

Initially, I investigated introducing derived temporal features, e.g., by discretizing the hour of the day into ‘day’ and ‘night’ intervals. However, as shown in fig. 3, the average fraction of available bikes throughout the day differs substantially between the training data and the data from the first ten stations (section 1). Hence, I considered that a feature of this kind would be unlikely to generalize well to the test data. Additionally, having elected to investigate tree models (section 4), which partition the spaces of quantitative features (Flach 2012, p. 155), discretization may be irrelevant.

4 Sub-task 1

4.1 Model classes

Gradient-boosted decision trees are a popular choice of model class for time-series forecasting problems (Bojer and Meldgaard 2021). In particular, `scikit-learn` provides an optimized implementation of gradient-boosted decision trees inspired by `LightGBM` (Ke et al. 2017). An advantage of this model class is that it supports missing values, which are evident in the data (table 2). For both cases of sub-task 1, I elected to compare the performance of individual decision trees and gradient-boosted decision trees. As a baseline, I predicted the arithmetic mean of the fraction of available bikes.

⁶See `sklearn.compose.TransformedTargetRegressor`.

Parameter	Values
criterion	<u>absolute_error</u>
max_depth	<u>None</u> , 1, 2, 5, 10, 20, 50
max_leaf_nodes	<u>None</u> , <u>7</u> , 15, 31, 63
min_samples_leaf	1, 2, 5, 10, <u>20</u> , 50, 100

(a) Decision tree

Parameter	Values
loss	<u>absolute_error</u>
scoring	<u>neg_mean_absolute_error</u>
l2_regularization	0.1, <u>0.2</u> , 0.5, 1.0
learning_rate	0.02, 0.05, <u>0.1</u> , 0.2, 0.5
max_depth	<u>None</u> , 1, 2, 5, 10, 20, <u>50</u> , 100
max_iter	20, 50, <u>100</u> , 200, 500
max_leaf_nodes	<u>None</u> , 7, <u>15</u> , 31, 63
min_samples_leaf	1, <u>2</u> , 5, 10, 20, 50, 100

(b) Gradient-boosted decision tree

Table 3: The parameter grids over which I performed hyperparameter search for sub-task 1. Except where stated, the default values were used, and the parameters of the best estimator are underlined. For a description of the parameters and their default values, see `sklearn.tree.DecisionTreeRegressor` and `sklearn.ensemble.HistGradientBoostingRegressor`.

The best estimator for each model class was determined by grid search with ten-fold cross-validation. With a separate model for each of the 75 stations, I found that it was too computationally expensive to perform grid search over a wide range of hyperparameters. Hence, I performed grid search for the case of a single model for all stations, then chose the values of the best estimator for the case of a separate model for each station. I fixed the scoring criteria to the mean absolute error, as per section 2. The resultant parameter values are listed in table 3.

4.2 Results

The mean scores achieved by the best estimator for each model class over the cross-validation folds and stations are listed in table 4. With a separate model for each station, there is greater variance in the mean score because there are 75 times as many cross-validation folds with commensurately fewer instances per fold. Intuitively, the baseline achieved a lower score with a separate model for each station. For both cases of sub-task 1, both model classes achieved a lower score than the baseline.

To determine whether the differences between the mean scores achieved by the best estimators were statistically significant, I performed paired t -tests of the scores on the cross-validation folds. The null hypothesis was that the differences between the mean scores achieved by the best estimators were due to chance. I did not compare the scores between estimators for the different cases of this sub-task because the samples are not paired. With separate models for each station, the best decision trees and gradient-boosted decision trees performed significantly better than the baseline for more than half of the stations (table 5a). The best gradient-boosted decision trees only performed significantly better than the best decision trees for seven stations. With a single model for all stations, the best gradient-boosted decision tree performed significantly better than the best decision tree (table 5b).

Model	μ	σ^2	Test
Baseline	4.43	3.82	4.47
Decision tree	3.53	3.52	3.01
Gradient-boosted decision tree	3.37	2.61	TODO

(a) Separate models for each station. μ is the mean of the scores for the 75 stations and ten cross-validation folds.

Model	μ	σ^2	Test
Baseline	5.45	0.06	TODO
Decision tree	3.51	0.08	TODO
Gradient-boosted decision tree	2.60	0.08	TODO

(b) Single model for all stations. μ is the mean of the scores for the ten cross-validation folds only.

Table 4: The mean scores and variances achieved by the best estimators for each model class for each case of sub-task 1 on the training data, and the corresponding score on the held-out test set (section 1).

Model 1	Model 2	$t > 0$	$p < 0.05$
Baseline	Decision tree	70	39
Baseline	Gradient-boosted decision tree	70	44
Decision tree	Gradient-boosted decision tree	44	7

(a) Separate models for each station. The numbers of positive t -statistics and significant p -values of the mean scores of the best estimators for each model class on the ten cross-validation folds for each of the 75 stations.

Model 1	Model 2	t -statistic	p -value
Baseline	Decision tree	18.6	1.70×10^{-9}
Baseline	Gradient-boosted decision tree	31.6	1.55×10^{-10}
Decision tree	Gradient-boosted decision tree	8.2	1.83×10^{-5}

(b) Single model for all stations. The t -statistics and p -values for paired t -tests of the mean scores of the best estimators for each model class on the ten cross-validation folds for all 75 stations.

Table 5: The results of paired t -tests of the scores achieved by the best estimators for each model class for each case of sub-task 1. A positive t -statistic indicates that ‘Model 2’ achieved a lower mean score than ‘Model 1’.

Model	temperature	bikes_3h	bikes_avg_full	bikes_avg_short	bikes_3h_diff_avg_full	bikes_3h_diff_avg_short
rlm_full		✓	✓		✓	
rlm_full_temp	✓	✓	✓		✓	
rlm_short		✓		✓		✓
rlm_short_full		✓	✓	✓	✓	✓
rlm_short_full_temp	✓	✓	✓	✓	✓	✓
rlm_short_temp	✓	✓		✓		✓

Table 6: The features used by the different kinds of linear models for sub-task 2. The features are listed in table 1 and follow the same ordering.

5 Sub-task 2

5.1 Model classes

The available information for sub-task 2 is a set of linear models that were trained on the number of available bikes at each of a separate set of 200 stations (section 1). Each of the six kinds of model uses a different set of features, which are listed in table 6. Unlike the gradient-boosted decision tree implementation in `scikit-learn`, these models do not support missing values, so I imputed them with the mean value of the feature over the data provided for sub-task 1. I used `sklearn.ensemble.StackingRegressor` to combine the predictions of the models by stacked generalization (Wolpert 1992). This model class is a meta-estimator that trains a second-level regressor on the predictions of a set of first-level regressors – in this case, the pre-trained models.

For this sub-task, I also excluded zero-variance and highly correlated features, but retained the ‘profile’ features (section 3.1). However, the `full` and `short` features are perfectly correlated, and the latter were excluded for sub-task 1. Therefore, I first assessed the performance of the different kinds of pre-trained models. A box plot of the mean scores achieved on the data provided for sub-task 1 is shown in fig. 4. It indicates that individual `rlm_short` and `rlm_short_temp` models generally achieved the lowest mean scores. However, with the default second-level regressor (ridge regression), the results of stacked generalization with each kind of model were very similar, so I elected to include all the models in the final analysis. For the second-level regressor, I compared ridge regression, decision trees, and gradient-boosted decision trees, as in section 4.1. The best second-level regressor was likewise determined by grid search with ten-fold cross-validation (section 2).

5.2 Results

The mean scores achieved by the best estimator for each model class over the cross-validation folds are listed in table 7. All the model classes achieved a lower score than the baselines in sub-task 1. As in section 4, I performed paired *t*-tests of the scores on the cross-validation folds to determine whether the differences between the mean scores achieved by the best estimators were statistically significant. The results are listed in table 8. Stacked generalization with both ridge regression and gradient-boosted decision trees as the second-level regressor achieved a lower score than with decision trees, but they did not perform significantly better than each other.




Figure 4: A box plot of the mean scores achieved by the different kinds of models for sub-task 2 over the ten cross-validation folds of the data from the 75 stations for the month of October 2014, i.e., the same data as for the second case of sub-task 1. The kinds of models are listed in table 6.

Model	μ	σ^2	Test
Ridge	2.85	0.13	TODO
Decision tree	3.79	0.10	TODO
Gradient-boosted decision tree	2.84	0.19	TODO

Table 7: The mean scores and variances achieved by the best estimators for each model class for sub-task 2 on the training data, and the corresponding score on the held-out test set (section 1).

Model 1	Model 2	t -statistic	p -value
Ridge	Decision tree	-20.7	6.70×10^{-9}
Ridge	Gradient-boosted decision tree	0.46	6.55×10^{-1}
Decision tree	Gradient-boosted decision tree	17.2	3.46×10^{-8}

Table 8: The results of paired t -tests of the scores achieved by the best estimators for each model class for sub-task 2. As in table 5, a positive t -statistic indicates that ‘Model 2’ achieved a lower mean score than ‘Model 1’.

References

- Alin, Aylin (2010). “Multicollinearity”. In: *WIREs Computational Statistics* 2.3, pp. 370–374.
- Bergmeir, Christoph, Rob J. Hyndman, and Bonsoo Koo (2018). “A note on the validity of cross-validation for evaluating autoregressive time series prediction”. In: *Computational Statistics & Data Analysis* 120, pp. 70–83.
- Bojer, Casper Solheim and Jens Peder Meldgaard (2021). “Kaggle forecasting competitions: An overlooked learning opportunity”. In: *International Journal of Forecasting* 37.2, pp. 587–603.
- Buitinck, Lars, Gilles Louppe, and Mathieu Blondel (2013). “API design for machine learning software: experiences from the scikit-learn project”. In: *ECML/PKDD 2013 Workshop: Languages for Data Mining and Machine Learning*.
- Flach, Peter (2012). *Machine Learning: The Art and Science of Algorithms that Make Sense of Data*. 1st ed. Cambridge University Press.
- Ke, Guolin et al. (2017). “LightGBM: A Highly Efficient Gradient Boosting Decision Tree”. In: *Advances in Neural Information Processing Systems*. Vol. 30. Curran Associates, Inc.
- Pedregosa, Fabian et al. (2011). “Scikit-learn: Machine Learning in Python”. In: *The Journal of Machine Learning Research* 12, pp. 2825–2830.
- Wolpert, David H. (1992). “Stacked generalization”. In: *Neural Networks* 5.2, pp. 241–259.